

# **Project Proposal:**

## **Simulating Buffer Manager Strategies for Join / Selection Queries**

**Group:** disks\_overloaded

**Members:**

- Utsav Basu (20CS30057)
- Anamitra Mukhopadhyay (20CS30064)
- Rounak Saha (20CS30043)
- Swarup Padhi (20CS30062)
- Spandan Halder (20CS10063)

**Introduction:**

This project aims to simulate a small buffer pool for simple join/selection queries on a few small tables. Popular buffer manager strategies such as LRU/MRU/CLOCK/Pinned blocks will be simulated, and these strategies will be compared in terms of the number of disk I/O required. For a more realistic simulation, we will use the SQLite C Library.

**Objective:**

The main objective of this project is to compare and analyse the performance of different buffer manager strategies for Join/Selection queries in terms of the number of disk I/O required.

**Methodology:**

The simulation will be implemented in the C/C++ language, and the SQLite C library will be used for more realistic simulation. We will create a small database with a few small tables that will be used for

join/selection queries. We will implement the following buffer manager strategies:

**Least Recently Used (LRU):** In this strategy, the block that has not been accessed for the longest time will be replaced.

**Most Recently Used (MRU):** In this strategy, the block that has been accessed most recently will be replaced.

**CLOCK:** In this strategy, the buffer manager maintains a circular list of blocks. The blocks are marked with a bit indicating whether they have been accessed since the last time they were considered. When a block is to be replaced, the buffer manager scans the list of blocks in a circular fashion, and the first block with the access bit set to 0 is replaced. If all the blocks have their access bit set to 1, the buffer manager starts again from the beginning of the list.

**Pinned Blocks:** In this strategy, some blocks are marked as pinned, which means they cannot be replaced.

We will simulate join/selection queries using each of these strategies and record the number of disk I/O operations required for each strategy.

### **Flowchart:**

Create a small database with a few small tables

Initialize buffer pool

### **Implement LRU strategy**

- a. Access the buffer pool
- b. Replace the least recently used block
- c. Record the number of disk I/O operations required

### **Implement MRU strategy**

- a. Access the buffer pool
- b. Replace the most recently used block
- c. Record the number of disk I/O operations required

### **Implement CLOCK strategy**

- a. Access the buffer pool
- b. Scan the list of blocks in a circular fashion
- c. Replace the first block with the access bit set to 0
- d. Record the number of disk I/O operations required

### **Implement Pinned Blocks strategy**

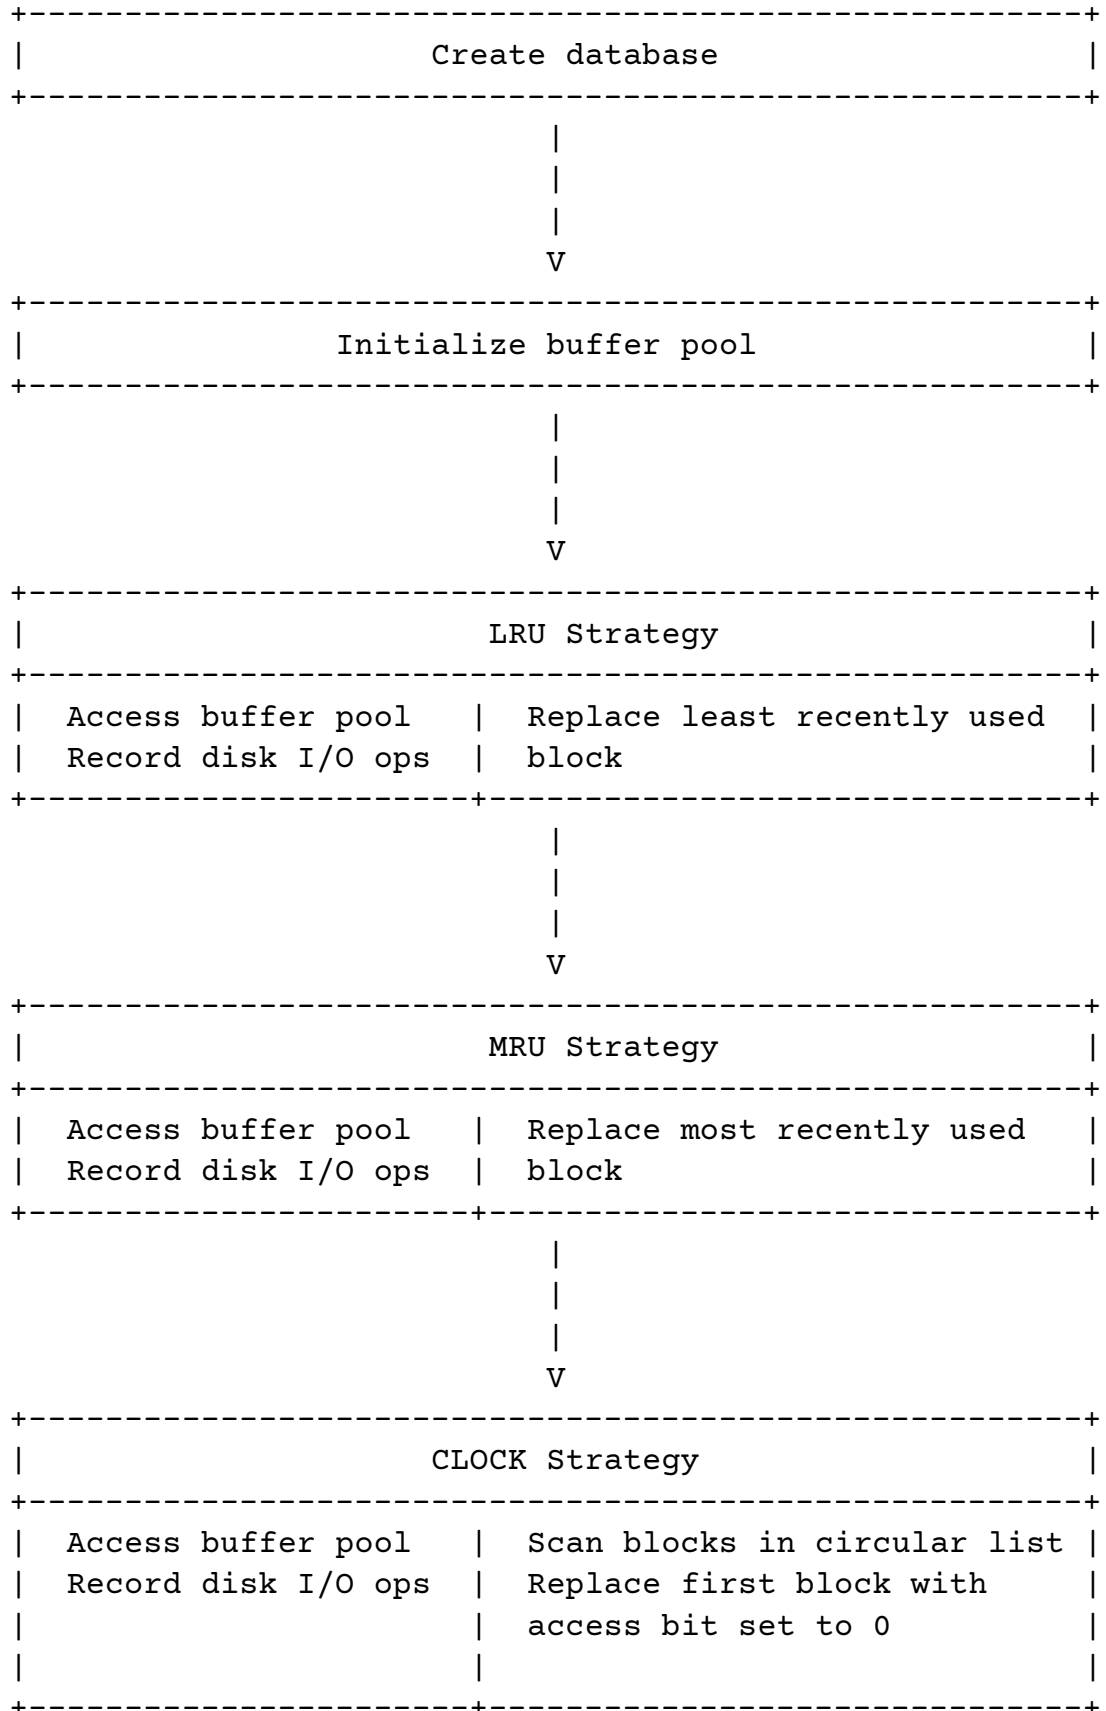
- a. Pin some blocks in the buffer pool
- b. Access the buffer pool
- c. Replace unpinned blocks
- d. Record the number of disk I/O operations required

Analyse and compare the performance of each strategy in terms of the number of disk I/O operations required.

Conclude and suggest the best strategy for join/selection queries.

### **Conclusion:**

In conclusion, this project aims to simulate a small buffer pool for simple join/selection queries on a few small tables. The popular buffer manager strategies such as LRU/MRU/CLOCK/Pinned blocks will be simulated, and their performance will be compared in terms of the number of disk I/O required. This project will provide insights into the best buffer manager strategy for join/selection queries.



|  
|  
|  
V

Pinned Blocks Strategy			
Pin blocks		Access buffer pool	
Access buffer pool		Replace unpinned blocks	
Record disk I/O ops			

|  
|  
|  
V

Analyze and compare performance			
---------------------------------	--	--	--

|  
|  
|  
V

Draw conclusions and suggest best strategy for queries			
---	--	--	--